

MACHINE LEARNING

Ludovico Bessi

July 2018

Contents

1	Week 1	1
1.1	Linear regression	1
2	Week 2	2
2.1	Multivariate linear regression	2
2.2	Alternative to Gradient descent: Normal equation	2

1 Week 1

1.1 Linear regression

Suppose we have this data: feet and price of some houses. We want to be able to predict the price of the next house knowing only the feet. The best way to do that is using linear regression:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

We want to find θ_i such that the line behaves well with the data. This is the same as minimizing the following function, called cost function:

$$J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x_i) - y_i)^2$$

We can minimize that function using "Gradient descent". We need to repeat the following operation until convergence:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

This is the same as saying:

$$\theta_j := \theta_j - \alpha * \frac{1}{n} \sum_{i=1}^n (h_{\theta}(x^{(i)} - y^{(i)})x_j^{(i)})$$

The intuition behind this is that at the minimum the second term goes to 0 as the derivative goes to 0. α is called the learning rate, and it is decided by the user.

2 Week 2

2.1 Multivariate linear regression

Suppose now that we have multiple information about houses. Then, we want to find a function like this:

$$h_{\theta}(x) = \sum_{i=1}^n \theta_i x_i$$

We can vectorize that function, obtaining just:

$$h_{\theta}(x) = \theta^T x$$

The gradient descent method does not change. However, there is a way to speed things up. We can normalize the θ_i to have them all in the same value range. It is helpful because we minimize the risk of oscillations between values. We just need to do the following:

$$x_i := \frac{x_i - \mu_i}{s_i}$$

Where μ_i is the average and s_i is the standard deviation.

2.2 Alternative to Gradient descent: Normal equation

Gradient descent is an iterative method, it is good even when we have many features. However it is still iterative, meaning that in principle we may spend a lot of time computing if we choose the wrong α .

We can then use the normal equation if n is not too big:

$$\theta = (X^T X)^{-1} X^T y$$

We need to be careful for two reasons: it may be the case that $X^T X$ is not invertible, because there are linearly dependent features, or too many of them. Moreover, calculating θ in this way costs $O(n^3)$.